

Company:

DeepReason.ai – a spinout from the University of Oxford – is building an innovative Knowledge Graph Platform which brings advanced reasoning capabilities on big data. DeepReason.ai offers a scalable engine for enterprises to manage their own knowledge graphs, which is deployed in the cloud or on-premise. The platform has strong commercial uptake and is used in verticals such as logistics, banking and manufacturing (Client base: Central Bank of Italy, Sberbank, Boeing, Meltwater, Miele).

You:

You are a motivated software engineer, keen to apply your expertise to challenging real-world applications and high-performance AI software systems. You will play an integral role in our team of experienced researchers and entrepreneurs, shaping the company in its early stage, with plenty of opportunities to grow.

Responsibilities:

- Implementing Knowledge Graph solutions to clients: developing data solutions with logical reasoning, ETL, scripting components (graph querying on large knowledge bases, building logical inference chains, data transformation pipelines). Technologies used: Datalog, Python, Jupyter, Apache Airflow, Amazon ECS, Postgres, RDF.
- Interaction with client engineering team. Capturing requirements, writing specifications, liaising with internal team
- Monitor and optimize ETL pipelines to serve regular data updates to client

Must-have:

- Strong degree (BSc or above) in Computer Science or related field, with good foundations in general computer science (data structures and algorithms).
- Minimum 1 years of experience in advanced SQL and Python scripting
- Experience in managing ETLs

- Familiarity with AWS, Git, Jupyter, Container environments (Docker)
- Ability to manage own workload and associated activities and interact well within a team environment. Plan and organise work independently to meet deadlines depending on programme and project priorities.
- Good communication and interpersonal skills. Very good professional English

Desirable:

- Experience in Logic programming (Prolog, Datalog)
- Familiarity with RDF stores, SPARQL
- Familiarity with CI/CD pipelines
- Fluency in Russian is a plus